Personal Agent - Usage Guide

A local-first assistant you run on your own computer. You chat with it in a web UI (or from Telegram), and it can read and summarize your mail, send mail you approve, read Telegram channels, fill out job applications from links, and help on LinkedIn - always with your confirmation before anything is sent.

1. Install and start

One command (macOS/Linux): `./start.sh` **Windows:** double-click `start.bat`
macOS Finder: double-click `start.command`

Requires Node.js 22 or newer from nodejs.org. The first run installs dependencies and a browser, then opens `http://127.0.0.1:8787` with the setup wizard. Everything - settings, keys, chat history, your profile - stays on your machine under the `data/` folder with private (600) permissions. Secrets are never written to logs.

2. Setup wizard

Step	What it asks	Where to get it
Model provider	Offline mock, or any OpenAI-compatible endpoint + key	platform.openai.com, OpenRouter, or local Ollama at http://localhost:11434/v1
Mail	Gmail address + app password	Google Account > Security > 2-Step Verification > App passwords
Telegram	Bot token; optional api_id/api_hash + channels	@BotFather in Telegram; my.telegram.org for api credentials
Profile	Name, email, phone, resume path	Used to answer application forms - never invented

Skip anything and finish later in **Settings**.

3. Everyday use

Mail: "check mail" - "unread mail" - "mail search stripe" - "summarize mail" - "send mail" (you review the exact email and confirm before it leaves).

Job applications: paste a job link with the word *apply*. The agent opens a real browser, fills the form from your saved profile, asks you in chat for anything missing (it never makes facts up), shows a summary, and submits only after you press **Confirm**. Every attempt - submitted, cancelled, failed - is logged under **Applications**.

Telegram: with a bot token saved, message your bot from Telegram - same agent, same chat. "telegram channels" reads the latest posts from channels you configure (one-time login in Settings).

LinkedIn (opt-in): enable in Settings after reading the risk note. "linkedin login" opens a browser so you log in once yourself. Then: "linkedin feed", "linkedin apply <url>", "linkedin message <url>: text", "linkedin invite <url>: note". Actions are paced (about one feed check an hour, daily caps) and stop immediately on any checkpoint or captcha.

4. Honest status labels

The sidebar shows what is actually working on your machine: green = verified on this install, amber = limited (offline mock model, or at-your-risk LinkedIn), grey = not set up. Features that need your real accounts (mail, Telegram, LinkedIn, applying) are only marked ready after a successful test from the Settings page.

LinkedIn risk note. LinkedIn's terms discourage automation. This feature uses your own logged-in session at a slow pace with daily caps, but your account could still be restricted. Use at your own risk.

5. Safety model

Thing	Rule
Sending mail	Only after you confirm the exact recipient, subject and body
Submitting applications	Only after a summary you confirm; cancelled on timeout
Form answers	Only from your saved profile or your chat answers - never invented
Secrets	Local file, mode 600, masked in logs, never uploaded
LinkedIn	Off by default, opt-in, paced, stops on checkpoints

6. Tests and development

`npm test` runs the offline test suite (intent parsing, secrets handling, the confirmation gates, the OpenAI-compatible provider against a stub, and a full Playwright application-form flow against a local fixture). `npm run doctor` checks your install. `npm run dev` starts with auto-reload.